

Simple Parallel Data Structures-4 Solutions continue

by

William Gropp and Ewing Lusk

Exercise: Using nonblocking operations

```
#include <stdio.h>
#include "mpi.h"

/* This example handles a 12 x 12 mesh, on 4 processors only. */
/* This using nonblocking operations */
#define maxn 12

int main( argc, argv )
int argc;
char **argv;
{
    int rank, value, size, errcnt, toterr, i, j;
    MPI_Request r[4];
    int nreq;
    MPI_Status statuses[4];
    double x[12][12];
    double xlocal[(12/4)+2][12];

    MPI_Init( &argc, &argv );

    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );

    if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );

    /* xlocal[][0] is lower ghostpoints, xlocal[][maxn+2] is upper */

    /* Fill the data as specified */
    for (i=1; i<=maxn/size; i++)
        for (j=0; j<maxn; j++)
            xlocal[i][j] = rank;
    for (j=0; j<maxn; j++) {
        xlocal[0][j] = -1;
        xlocal[maxn/size+1][j] = -1;
    }

    /* Send up unless I'm at the top, then receive from below */
    /* Note the use of xlocal[i] for &xlocal[i][0] */
    nreq = 0;
```

```

    if (rank < size - 1)
        MPI_Isend( xlocal[maxn/size], maxn, MPI_DOUBLE, rank + 1, 0,
                   MPI_COMM_WORLD, &r[nreq++] );
    if (rank > 0)
        MPI_Irecv( xlocal[0], maxn, MPI_DOUBLE, rank - 1, 0,
MPI_COMM_WORLD,
                   &r[nreq++] );
    /* Send down unless I'm at the bottom */
    if (rank > 0)
        MPI_Isend( xlocal[1], maxn, MPI_DOUBLE, rank - 1, 1,
MPI_COMM_WORLD,
                   &r[nreq++] );
    if (rank < size - 1)
        MPI_Irecv( xlocal[maxn/size+1], maxn, MPI_DOUBLE, rank + 1, 1,
MPI_COMM_WORLD, &r[nreq++] );

    MPI_Waitall( nreq, r, statuses );

    /* Check that we have the correct results */
    errcnt = 0;
    for (i=1; i<=maxn/size; i++)
        for (j=0; j<maxn; j++)
            if (xlocal[i][j] != rank) errcnt++;
    for (j=0; j<maxn; j++) {
        if (xlocal[0][j] != rank - 1) errcnt++;
        if (rank < size-1 && xlocal[maxn/size+1][j] != rank + 1) errcnt++;
    }

    MPI_Reduce( &errcnt, &toterr, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD );
    if (rank == 0) {
        if (toterr)
            printf( "! found %d errors\n", toterr );
        else
            printf( "No errors\n" );
    }

    MPI_Finalize( );
    return 0;
}

```

Output

```

% mpicc -o exchng exchng.c
% mpirun -np 4 exchng
No errors

```

Makefile

```

# Generated automatically from Makefile.in by configure.
ALL: exchng
exchng: exchng.c
        mpicc -o exchng exchng.c

```

Exercise: Shifting data around

```
#include <stdio.h>
#include "mpi.h"

/* This example handles a 12 x 12 mesh, on 4 processors only. */
#define maxn 12

int main( argc, argv )
int argc;
char **argv;
{
    int rank, value, size, errcnt, toterr, i, j;
    int up_nbr, down_nbr;
    MPI_Status status;
    double x[12][12];
    double xlocal[(12/4)+2][12];

    MPI_Init( &argc, &argv );

    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );

    if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );

    /* xlocal[][0] is lower ghostpoints, xlocal[][maxn+2] is upper */

    /* Fill the data as specified */
    for (i=1; i<=maxn/size; i++)
        for (j=0; j<maxn; j++)
            xlocal[i][j] = rank;
    for (j=0; j<maxn; j++) {
        xlocal[0][j] = -1;
        xlocal[maxn/size+1][j] = -1;
    }

    /* Send up and receive from below (shift up) */
    /* Note the use of xlocal[i] for &xlocal[i][0] */
    /* Note that we use MPI_PROC_NULL to remove the if statements that
       would be needed without MPI_PROC_NULL */
    up_nbr = rank + 1;
    if (up_nbr >= size) up_nbr = MPI_PROC_NULL;
    down_nbr = rank - 1;
    if (down_nbr < 0) down_nbr = MPI_PROC_NULL;
```

```

MPI_Sendrecv( xlocal[maxn/size], maxn, MPI_DOUBLE, up_nbr, 0,
              xlocal[0], maxn, MPI_DOUBLE, down_nbr, 0,
              MPI_COMM_WORLD, &status );

/* Send down and receive from above (shift down) */
MPI_Sendrecv( xlocal[1], maxn, MPI_DOUBLE, down_nbr, 1,
              xlocal[maxn/size+1], maxn, MPI_DOUBLE, up_nbr, 1,
              MPI_COMM_WORLD, &status );

/* Check that we have the correct results */
errcnt = 0;
for (i=1; i<=maxn/size; i++)
    for (j=0; j<maxn; j++)
        if (xlocal[i][j] != rank) errcnt++;
for (j=0; j<maxn; j++) {
    if (xlocal[0][j] != rank - 1) errcnt++;
    if (rank < size-1 && xlocal[maxn/size+1][j] != rank + 1) errcnt++;
}

MPI_Reduce( &errcnt, &toterr, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD );
if (rank == 0) {
    if (toterr)
        printf( "! found %d errors\n", toterr );
    else
        printf( "No errors\n" );
}

MPI_Finalize( );
return 0;
}

```

Output

```

% mpicc -o exchng exchng.c
% mpirun -np 4 exchng
No errors

```

Makefile

```

# Generated automatically from Makefile.in by configure.
ALL: exchng
exchng: exchng.c
        mpicc -o exchng exchng.c

```

Exercise: Exchanging data with MPI_Sendrecv

```

#include <stdio.h>
#include "mpi.h"

```

```

/* This example handles a 12 x 12 mesh, on 4 processors only. */
#define maxn 12

int main( argc, argv )
int argc;
char **argv;
{
    int rank, value, size, errcnt, toterr, i, j;
    int up_nbr, down_nbr;
    MPI_Status status;
    double x[12][12];
    double xlocal[(12/4)+2][12];

    MPI_Init( &argc, &argv );

    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );

    if (size != 4) MPI_Abort( MPI_COMM_WORLD, 1 );

    /* xlocal[][0] is lower ghostpoints, xlocal[][maxn+2] is upper */

    /* Fill the data as specified */
    for (i=1; i<=maxn/size; i++)
        for (j=0; j<maxn; j++)
            xlocal[i][j] = rank;
    for (j=0; j<maxn; j++) {
        xlocal[0][j] = -1;
        xlocal[maxn/size+1][j] = -1;
    }

    /* Processors 0 and 1 exchange, 2 and 3 exchange, etc. Then
       1 and 2 exchange, 3 and 4, etc. The formula for this is
       if (even) exchnge up else down
       if (odd)  exchnge up else down
    */
    /* Note the use of xlocal[i] for &xlocal[i][0] */
    /* Note that we use MPI_PROC_NULL to remove the if statements that
       would be needed without MPI_PROC_NULL */
    up_nbr = rank + 1;
    if (up_nbr >= size) up_nbr = MPI_PROC_NULL;
    down_nbr = rank - 1;
    if (down_nbr < 0) down_nbr = MPI_PROC_NULL;

    if ((rank % 2) == 0) {
        /* exchange up */
        MPI_Sendrecv( xlocal[maxn/size], maxn, MPI_DOUBLE, up_nbr, 0,
                      xlocal[maxn/size+1], maxn, MPI_DOUBLE, up_nbr, 0,
                      MPI_COMM_WORLD, &status );
    }
}

```

```

    }
    else {
        /* exchange down */
        MPI_Sendrecv( xlocal[1], maxn, MPI_DOUBLE, down_nbr, 0,
                      xlocal[0], maxn, MPI_DOUBLE, down_nbr, 0,
                      MPI_COMM_WORLD, &status );
    }

    /* Do the second set of exchanges */
    if ((rank % 2) == 1) {
        /* exchange up */
        MPI_Sendrecv( xlocal[maxn/size], maxn, MPI_DOUBLE, up_nbr, 1,
                      xlocal[maxn/size+1], maxn, MPI_DOUBLE, up_nbr, 1,
                      MPI_COMM_WORLD, &status );
    }
    else {
        /* exchange down */
        MPI_Sendrecv( xlocal[1], maxn, MPI_DOUBLE, down_nbr, 1,
                      xlocal[0], maxn, MPI_DOUBLE, down_nbr, 1,
                      MPI_COMM_WORLD, &status );
    }

    /* Check that we have the correct results */
    errcnt = 0;
    for (i=1; i<=maxn/size; i++)
        for (j=0; j<maxn; j++)
            if (xlocal[i][j] != rank) errcnt++;
    for (j=0; j<maxn; j++) {
        if (xlocal[0][j] != rank - 1) errcnt++;
        if (rank < size-1 && xlocal[maxn/size+1][j] != rank + 1) errcnt++;
    }

    MPI_Reduce( &errcnt, &toterr, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD );
    if (rank == 0) {
        if (toterr)
            printf( "! found %d errors\n", toterr );
        else
            printf( "No errors\n" );
    }

    MPI_Finalize( );
    return 0;
}

```

Output

```

% mpicc -o exchng exchng.c
% mpirun -np 4 exchng
No errors

```

Makefile

Generated automatically from Makefile.in by configure.

ALL: exchnng

exchnng: exchnng.c

mpicc -o exchnng exchnng.c